



UNIVERSITÀ
DEGLI STUDI
FIRENZE

Mean Shift

Parallel Computing Final-Term

Angelo D'Amante
Fabian Greavu



Introduction

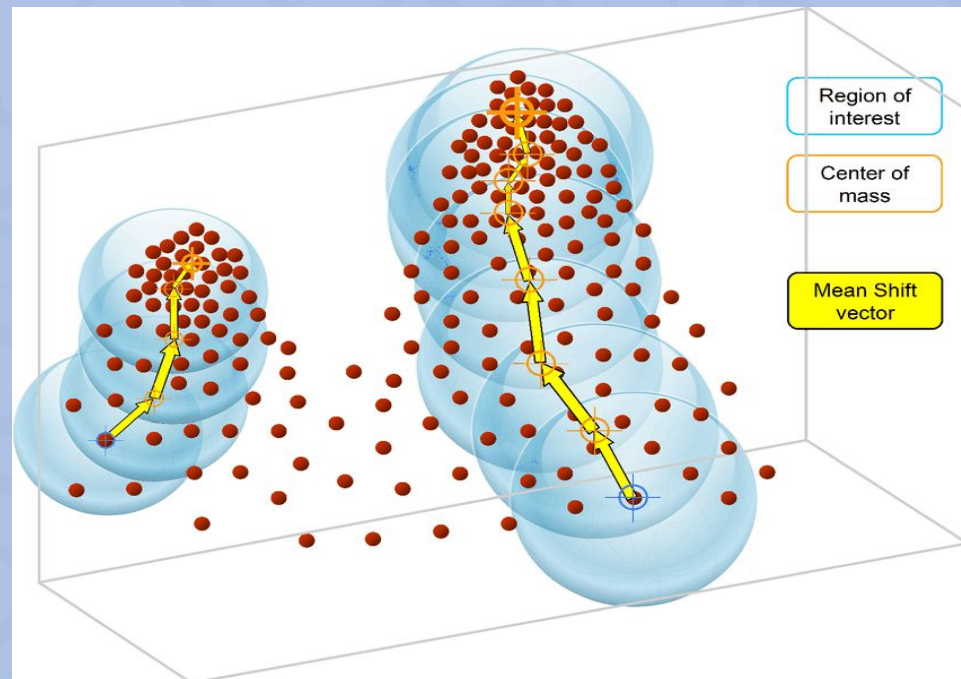
Introduction: Mean shift idea

Mean Shift is a non-parametric feature-space analysis technique for locating the maximas of a density function

Applications:

- clustering
- computer vision
- image processing

Main idea: move each point to the center of mass of his surroundings and iterate on the dataset until convergence.



Introduction: Mean shift parallelization

Mean shift algorithm has 2 main parts:

1. for 1, 2, ..., K:
 shift $x \leftarrow m(x)$ for all x in dataset
2. reduce dataset to centroids

Idea:
Step 1 Sequential \rightarrow Parallel

Introduction: Mean shift sequential algorithm

Pseudo-Code Algorithm:

```
new_data = original_points
while iteration < I do
  for p in new_data do
    num, den = 0
    for op in original_points do
      d = Distance(p, op)
      gaussian = K(dist)
      num = num + op * gaussian
      den = den + gaussian
    end for
    p = num/den
  end for
end while
return reduce_centroids(data, eps)
```

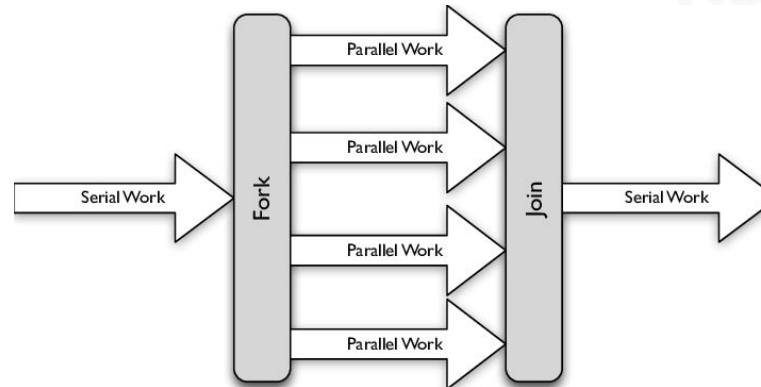
$$m(x) = \frac{\sum_{x_i \in N(x)} K(x_i - x)x_i}{\sum_{x_i \in N(x)} K(x_i - x)}$$

$$k(x) = e^{-\frac{x^2}{2\sigma^2}},$$



OpenMP Implementation

OpenMP: Introduction



- Advantages of the ***fork-join*** model.
 - Only few directives needed for this task.
 - Benefit from both ***shared*** and ***private*** memory.
-
- Exploit both workload sharing mechanism:
 - ***Static*** → ***set by User***
 - ***Dynamic*** → ***set by OpenMP***


```
#pragma omp parallel for default(none) shared(...) schedule(static) num_threads(num_thread)
```

```
1: procedure MEANSHIFT(original_points)
2:   new_data = original_points
3:   while iteration < I do
4:     for p in new_data do
5:       num = 0
6:       den = 0
7:       for op in original_points do
8:         d = Distance(p, op)
9:         gaussian = K(dist)
10:        num = num + op * gaussian
11:        den = den + gaussian
12:      end for
13:      p = num/den
14:    end for
15:  end while
16:  data = reduce_centroids(data,  $\epsilon$ )
17:  return data
```

Static Version

Dynamic Version

```
#pragma omp parallel for default(none) shared(...) schedule(dynamic)
```


Two CPU's architectures examined.

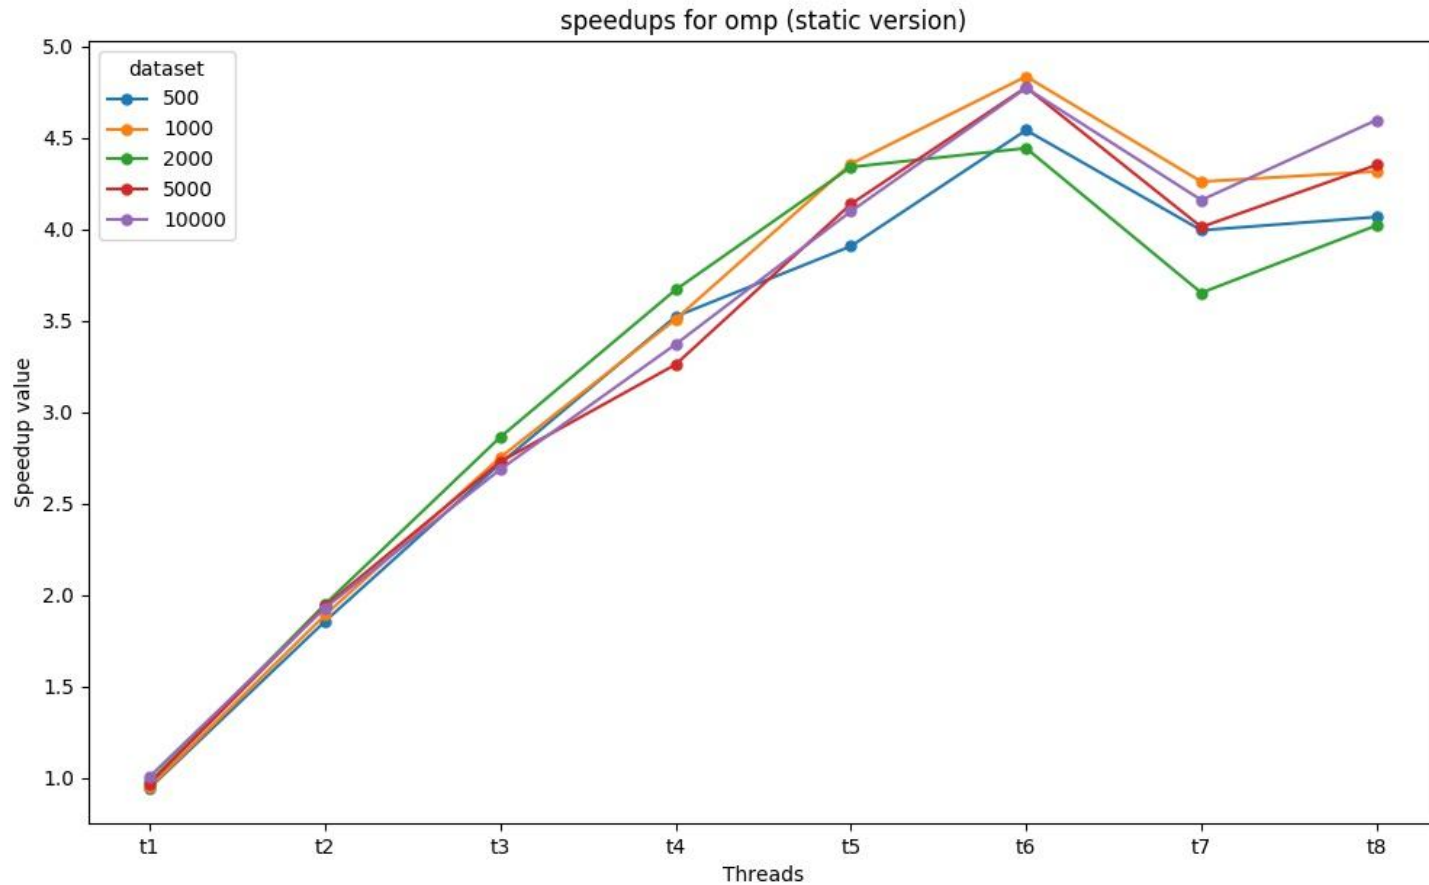
1. **Sequential** and **Parallel** timings are taken by:
CPU Intel(R) Core(TM) i7-8750H CPU @ 2.2GHz
2. **Sequential** and **Parallel** timings are taken by:
CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz
Each core boosts up to 4.6 Ghz OOTB



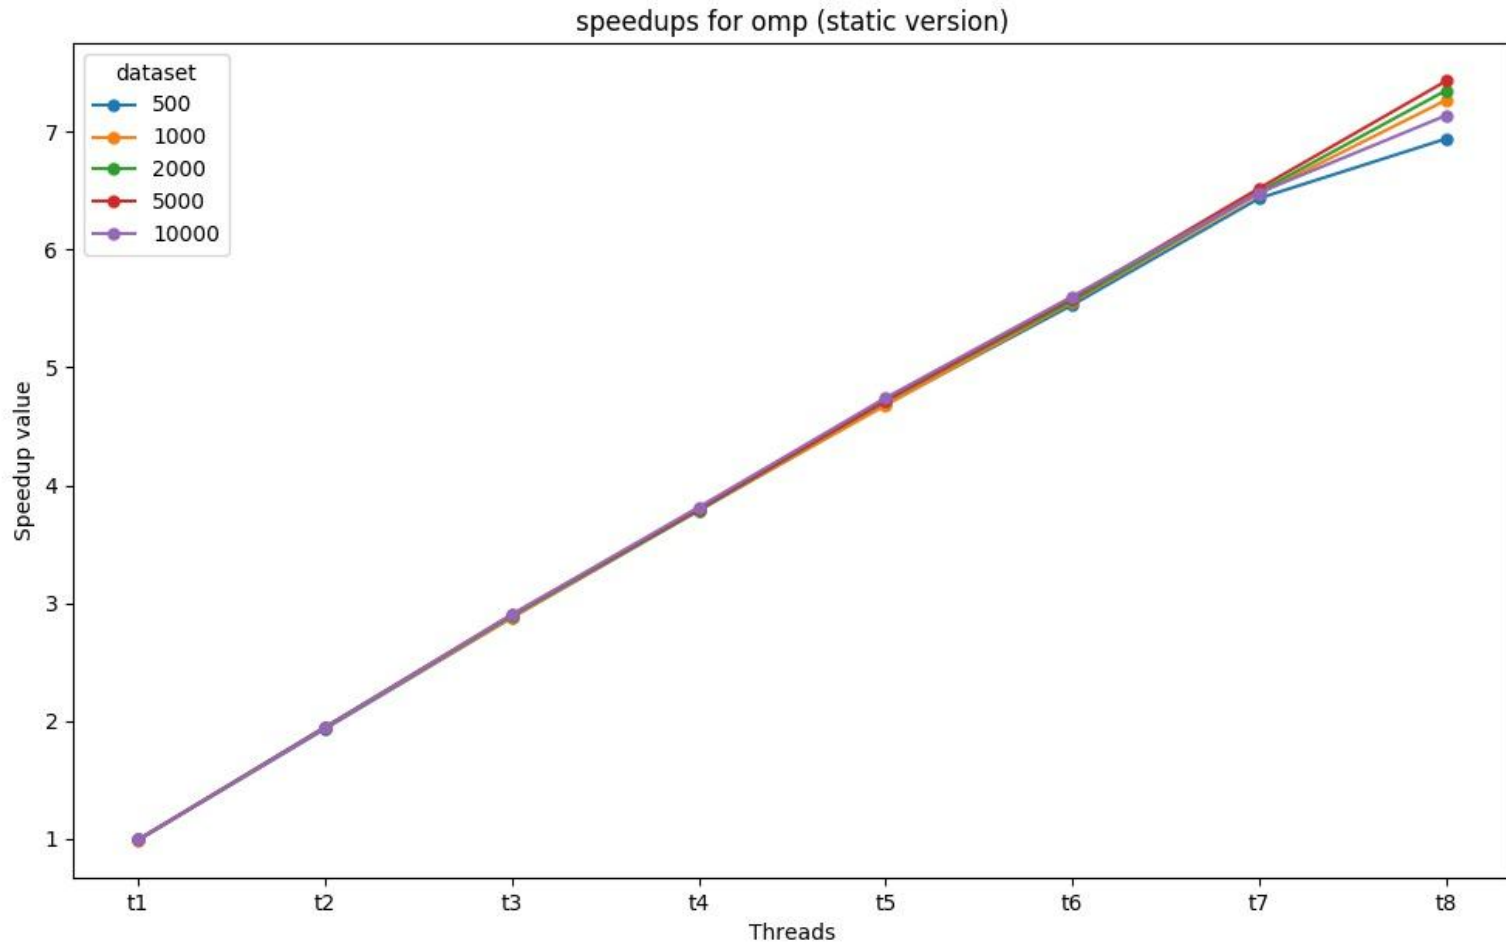
Dataset = {500, 1000, 2000, 5000, 10000}

Versions:

- Static with number of threads in [1, 8]
- Dynamic Version

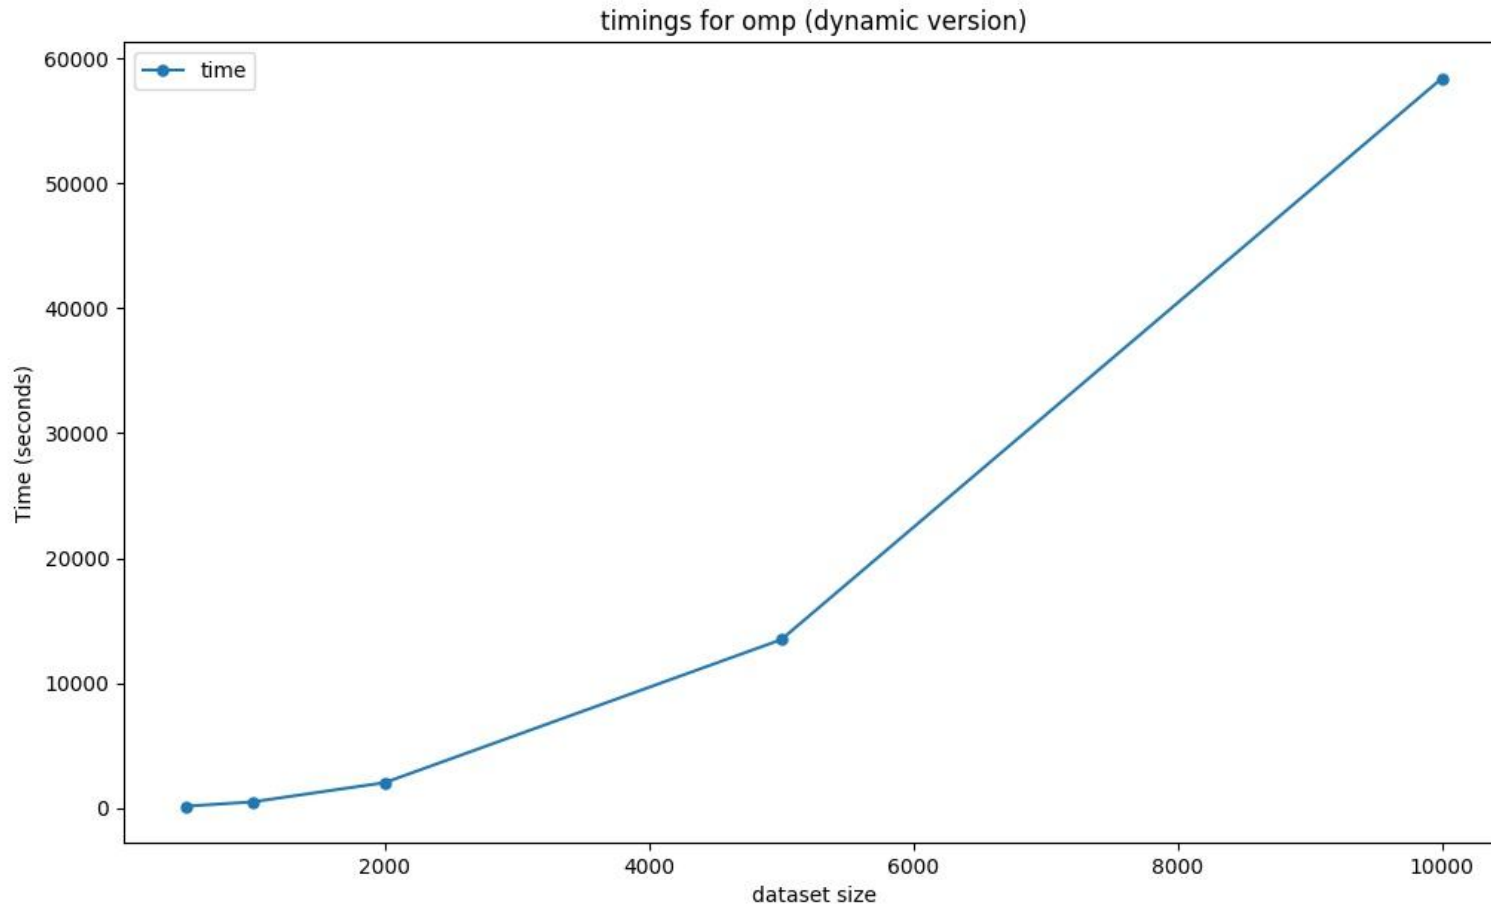


CPU Intel(R) Core(TM) i7-8750H CPU @ 2.2GHz

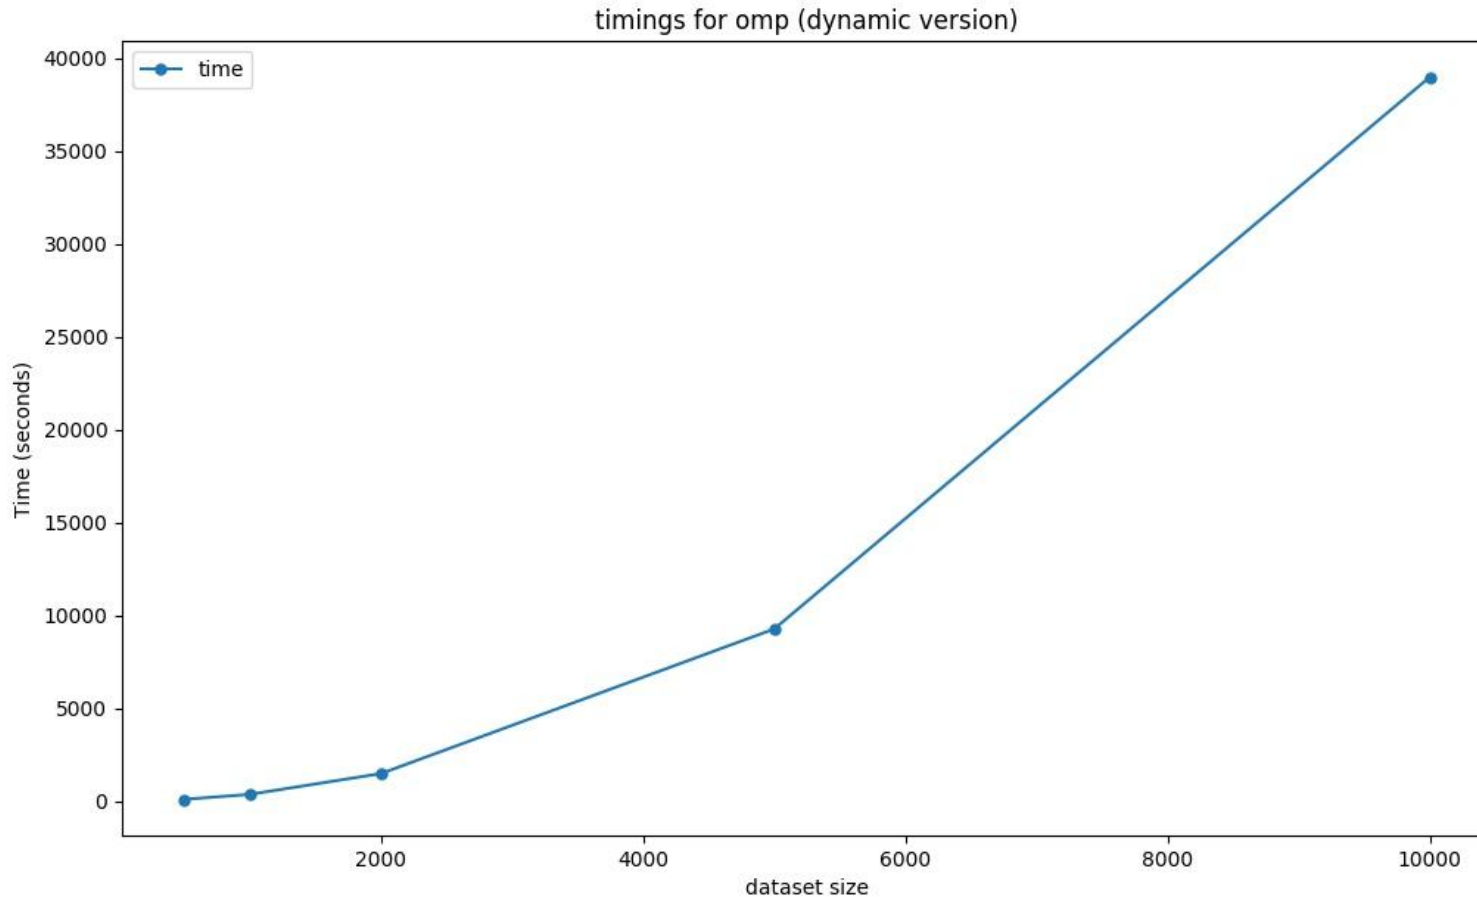


CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz

OpenMP: Results

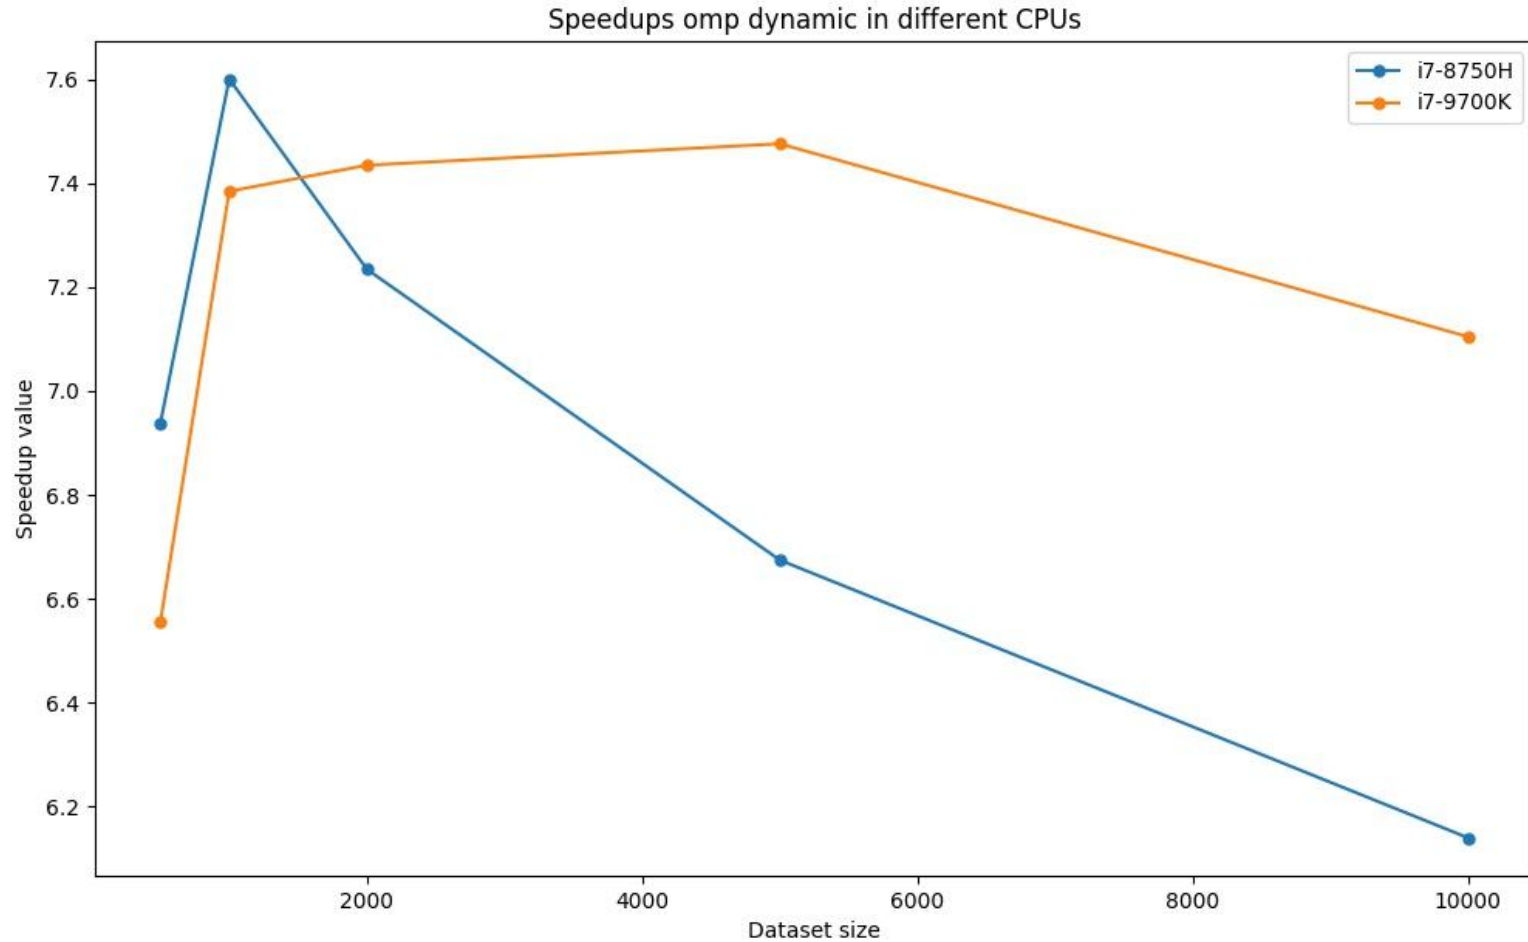


CPU Intel(R) Core(TM) i7-8750H CPU @ 2.2GHz



CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz

OpenMP: Results

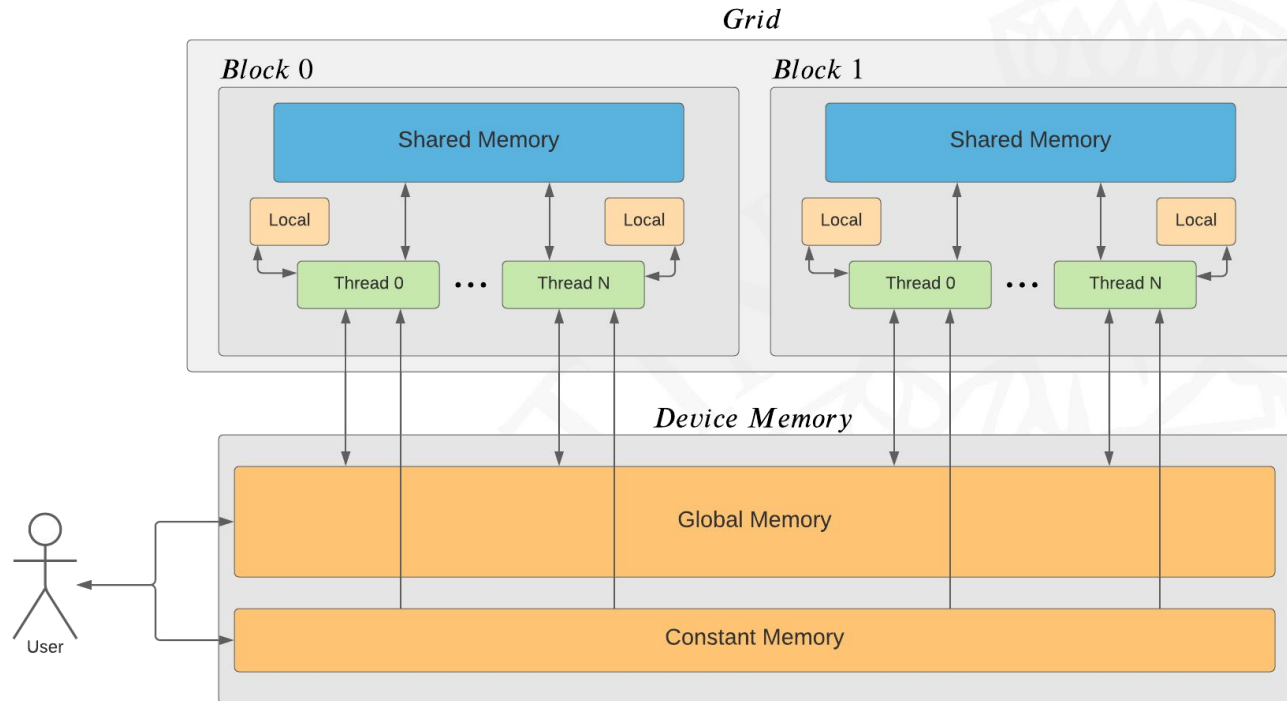


CPU Intel(R) Core(TM) i7-8750H CPU @ 2.2GHz vs CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz



CUDA Implementation

CUDA: Introduction



- **SIMT** (Single Instruction Multiple Thread) model.
- User decides number of threads $T \rightarrow \left\lceil \frac{N}{T} \right\rceil$ blocks.
- Exploit both **global** and **shared** memory access.

Naive:

```
csv = load_csv( ... );

float *dev_data, *dev_data_tmp;
size_t data_bytes = POINTS_NUMBER * D * sizeof(float);
cudaMalloc(&dev_data, data_bytes);
cudaMalloc(&dev_data_tmp, data_bytes);

cudaMemcpy(dev_data, data.data(), data_bytes,
cudaMemcpyHostToDevice);
cudaMemcpy(dev_data_tmp, data_next.data(), data_bytes,
cudaMemcpyHostToDevice);

for (size_t i = 0; i < NUM_ITER; ++i) {
    compute_weights_naive_kernel<<<BLOCKS,
    TW>>>(dev_data, dev_data_tmp, POINTS_NUMBER);
    cudaDeviceSynchronize();
    temp_data = dev_data;
    dev_data = dev_data_tmp;
    dev_data_tmp = temp_data;
}
cudaMemcpy(data.data(), dev_data, data_bytes,
cudaMemcpyDeviceToHost);
centers = resuce_to_centroids(data, MIN_DISTANCE)
```

```
const int TW = 64;
const int BLOCKS = (POINTS_NUMBER + TW - 1) / TW;
```

Shared Memory:

```
csv = load_csv( ... );

float *dev_data, *dev_data_tmp;
size_t data_bytes = POINTS_NUMBER * D * sizeof(float);
cudaMalloc(&dev_data, data_bytes);
cudaMalloc(&dev_data_tmp, data_bytes);

cudaMemcpy(dev_data, data.data(), data_bytes,
cudaMemcpyHostToDevice);
cudaMemcpy(dev_data_tmp, data_next.data(), data_bytes,
cudaMemcpyHostToDevice);

for (size_t i = 0; i < NUM_ITER; ++i) {
    compute_weights_shared_mem_kernel<<<BLOCKS,
    TW>>>(dev_data, dev_data_tmp, POINTS_NUMBER,
    BLOCKS);
    cudaDeviceSynchronize();
    temp_data = dev_data;
    dev_data = dev_data_tmp;
    dev_data_tmp = temp_data;
}
cudaMemcpy(data.data(), dev_data, data_bytes,
cudaMemcpyDeviceToHost);
centers = resuce_to_centroids(data, MIN_DISTANCE)
```

Naive:

```
__global__ void naive_kernel(data, data_tmp,
POINTS_NUMBER){
    int tid = (blockIdx.x * blockDim.x) + threadIdx.x;
    int x = tid*2, y = tid*2+1;
    if (tid < POINTS_NUMBER) {
        float new_position = {0., 0.}; // and other variables
        for (int i = 0; i < POINTS_NUMBER; ++i) {
            int xloop = i*2, yloop = i*2+1;
            eucl_dist = // calculate distance data[x], data[xloop]
            if (eucl_dist <= RADIUS) {
                weight = expf(-eucl_dist / BANDWIDTH);
                new_position += ... //update position
            }
        }
        data_tmp[x, y] = new_position // move p
    }
}
```

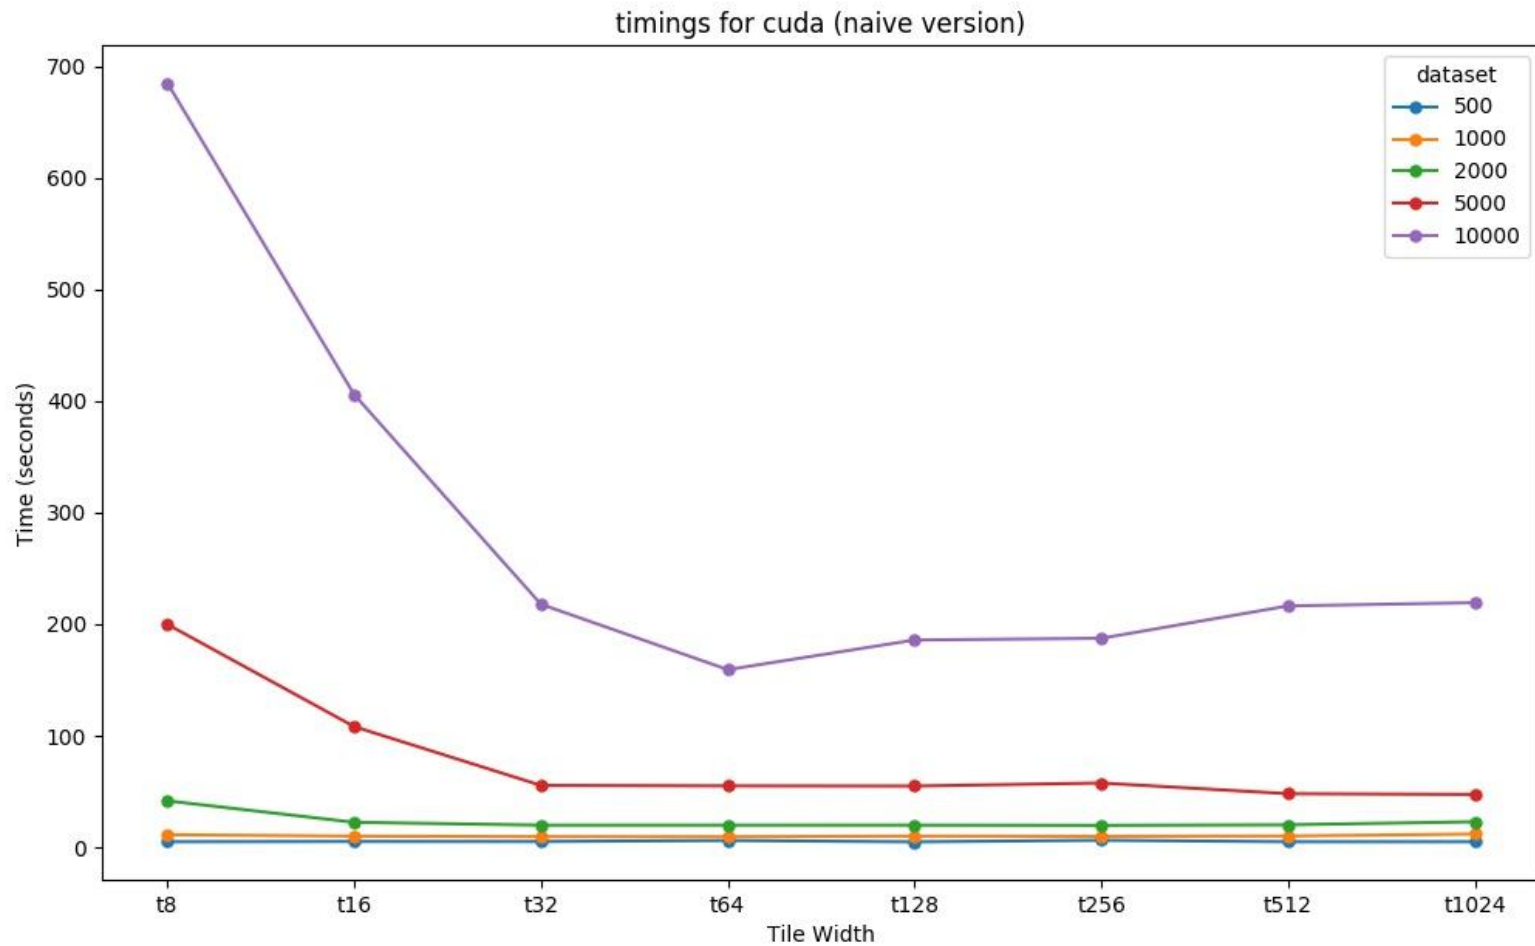
Shared Memory:

```
const int BLOCKS = (POINTS_NUMBER + TW - 1) / TW;
__global__ void sm_kernel(data, data_tmp, POINTS_NUMBER){
    __shared__ float local_data[TW * 2];
    __shared__ float flag_data[TW];
    int tid = (blockIdx.x * blockDim.x) + threadIdx.x;
    int x = tid*2, y = tid*2+1; // and other variables
    for (int t = 0; t < BLOCKS; ++t) {
        tid_in_tile = t * TILE_WIDTH + threadIdx.x;
        if (tid_in_tile < POINTS_NUMBER) {
            local_data[threadIdx.x * 2] = data[tid_in_tile * 2]
            local_data[threadIdx.x * 2 + 1] = data[tid_in_tile * 2 + 1]
            flag_data[threadIdx.x] = 1;
        }else{
            flag_data[threadIdx.x] = 0;
        }
        __syncthreads();
        for (int i = 0; i < TW; ++i) {
            int local_x_tile = i*2, local_y_tile = i*2+1;
            eucl_dist = ... // calculate distance dtat[x], data[x_in_tile]
            if (eucl_dist <= RADIUS) {
                weight = expf(-eucl_dist / BANDWIDTH);
                new_position += ... //update position data[x], data[local_x_tile]
            }
        }
        __syncthreads();
    }
    if (tid < POINTS_NUMBER) {new_position = ... // move p }
}
```

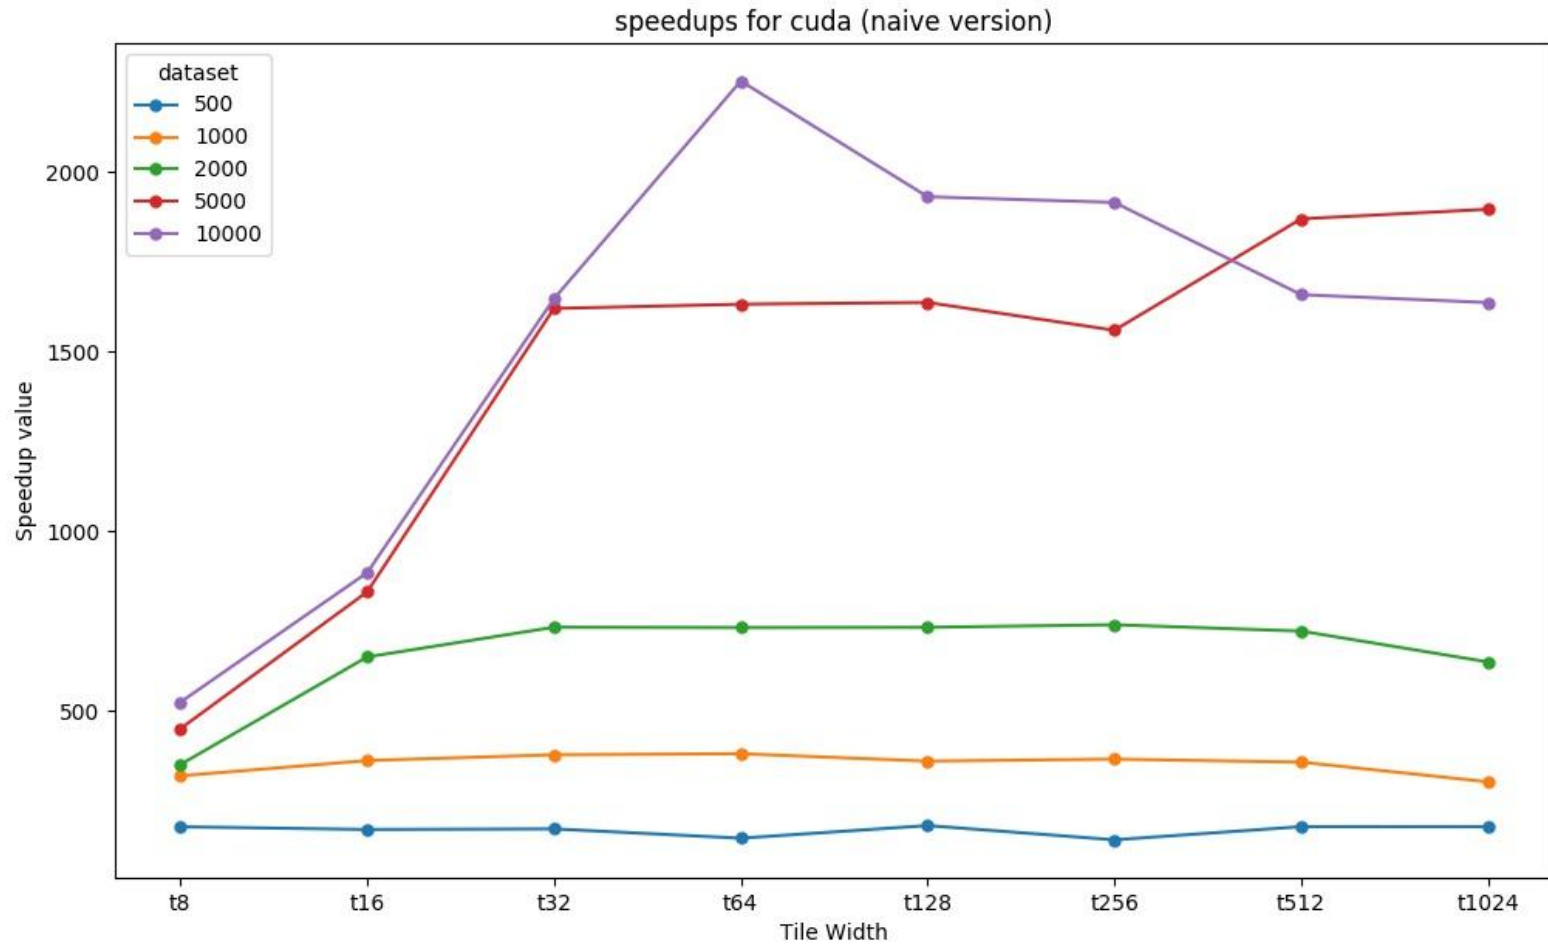
CUDA: Experiments

- **Sequential Timings** are taken by:
CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz
(Boosting at 4.6Ghz ootb)
- **Parallel Timings** are taken by:
GPU NVIDIA GeForce GTX 1050 Ti
With 4096 MB dedicated and 768 CUDA Cores
- **Experiments** with:
 - Dataset = {500, 1000, 2000, 5000, 10000}
 - Tile Width = {8, 16, 32, 64, 128, 256, 512, 1024}

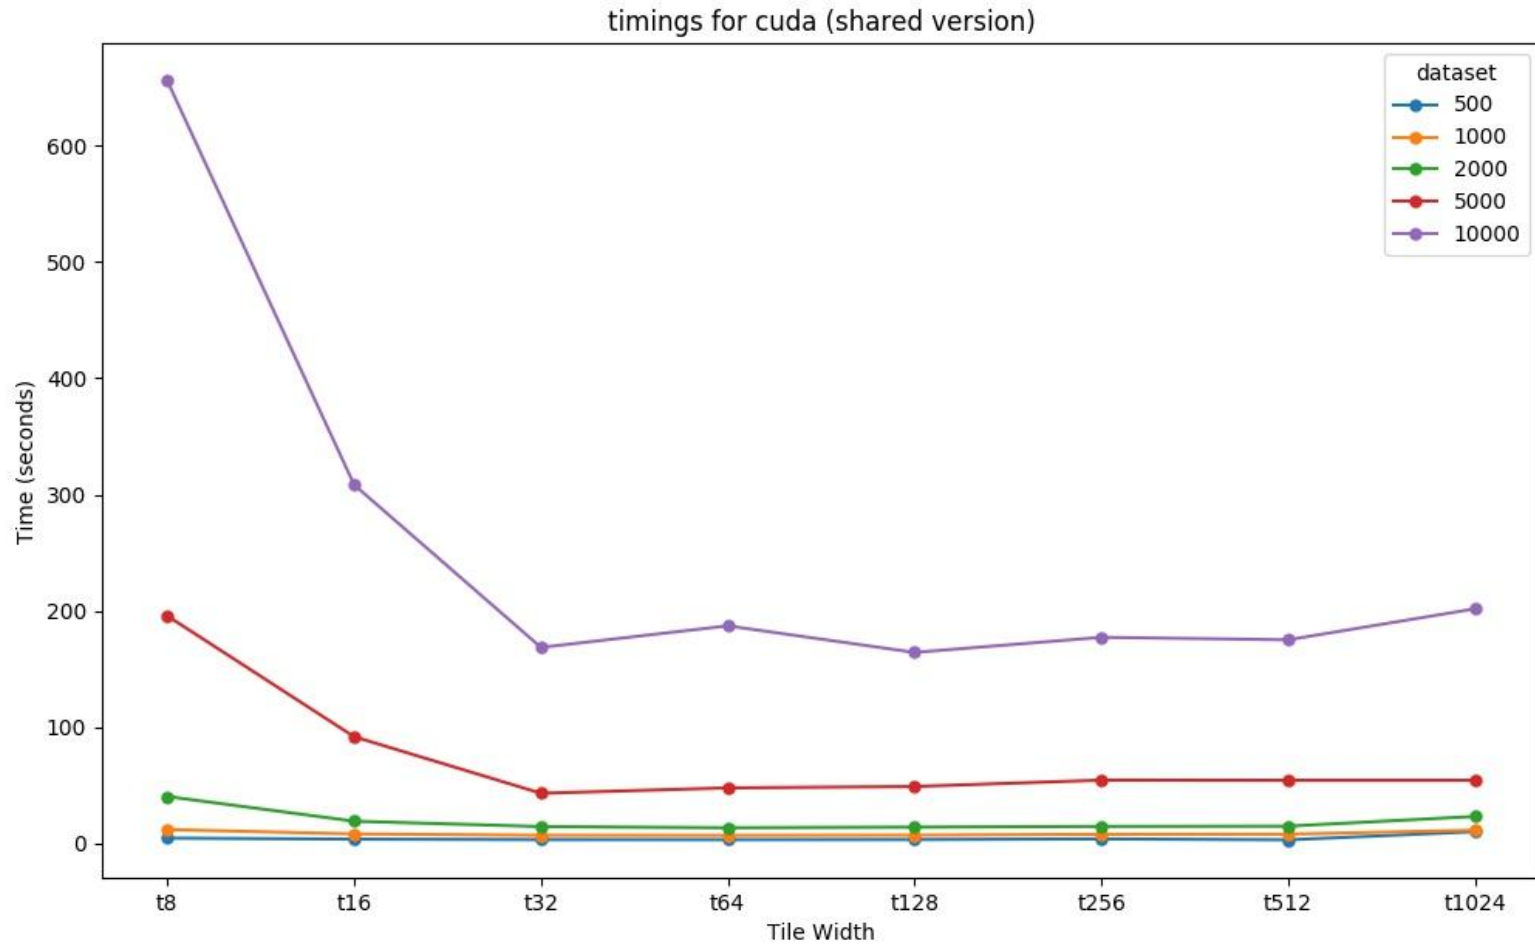
CUDA: naive results (timings)



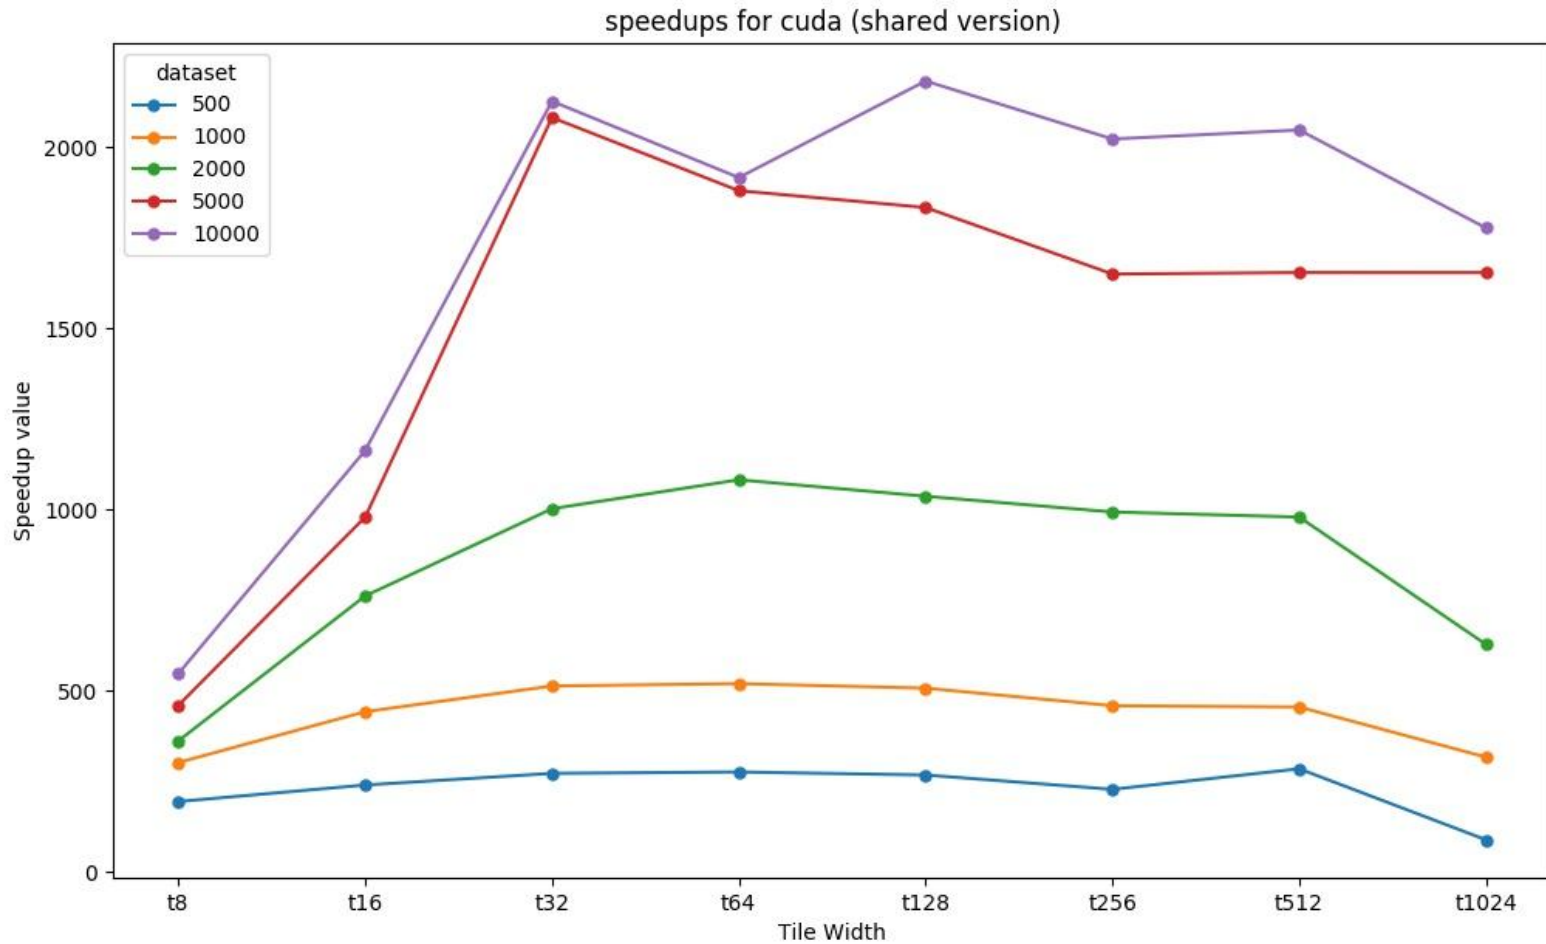
CUDA: naive results (speedups)



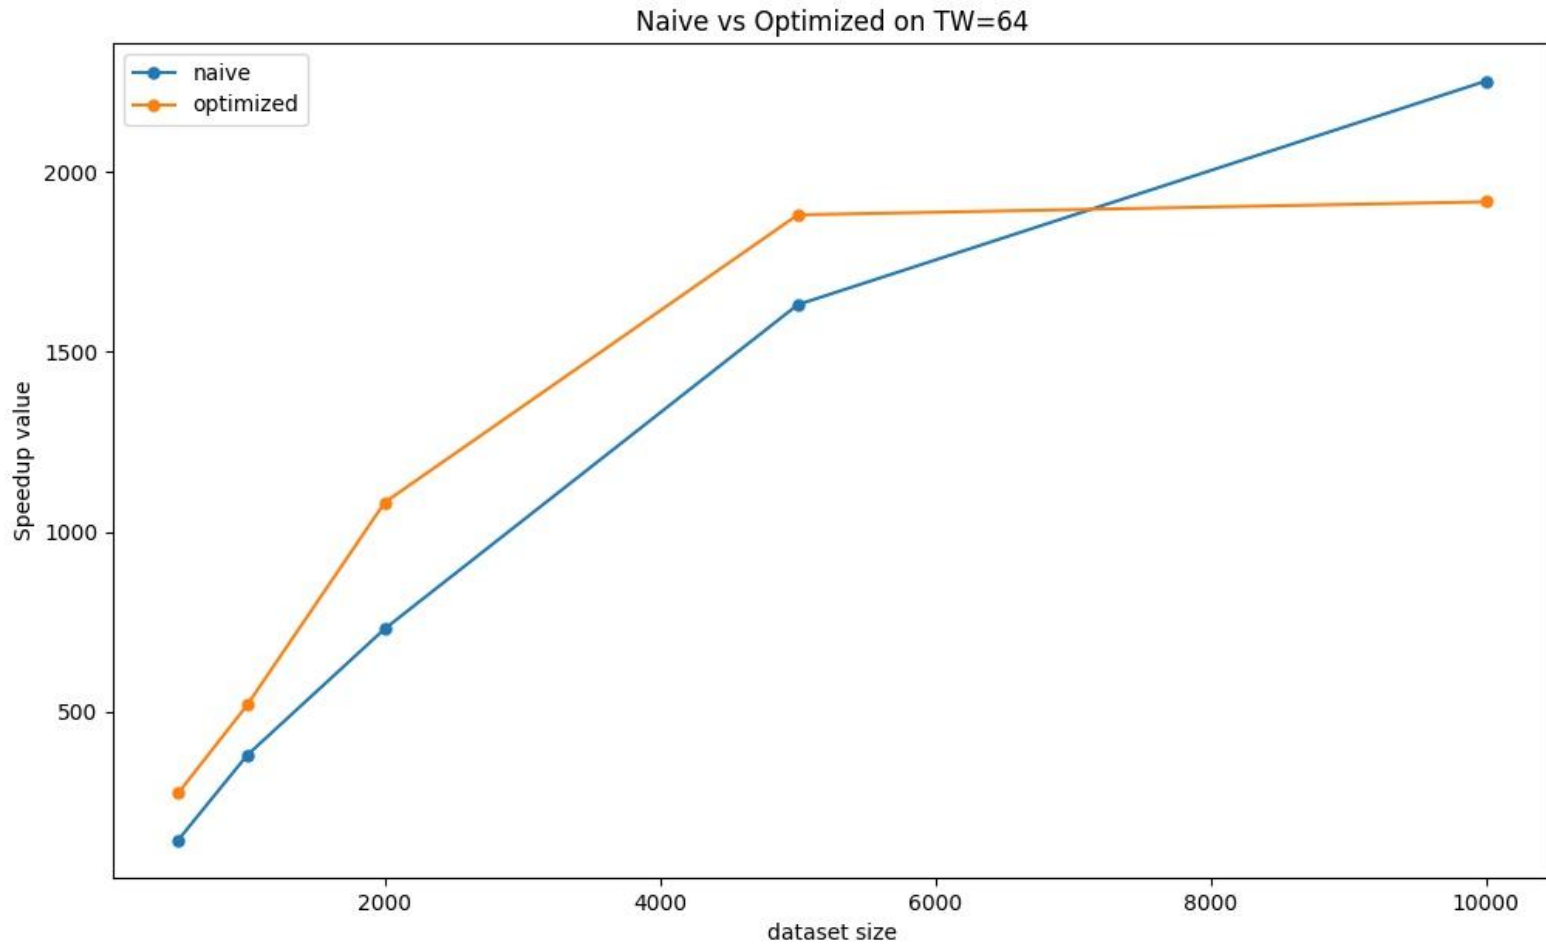
CUDA: shared memory results (timings)



CUDA: shared memory results (speedups)



CUDA: shared vs naive (speedups)



CUDA: comparison results (speedups)

Extra results: Comparison between different GPUS and architectures

- **Sequential Timings** are taken by:

CPU Intel(R) Core(TM) i7-9700K CPU @ 3.6GHz

- **Parallel Timings** on:

GPU NVIDIA GeForce **GTX 1050 Ti** m

Pascal

4096 MB dedicated and 768 CUDA Cores

GPU NVIDIA GeForce **GTX 1660 Ti**

Turing

6144 MB dedicated and 1536 CUDA Cores

GPU NVIDIA GeForce **RTX 2060**

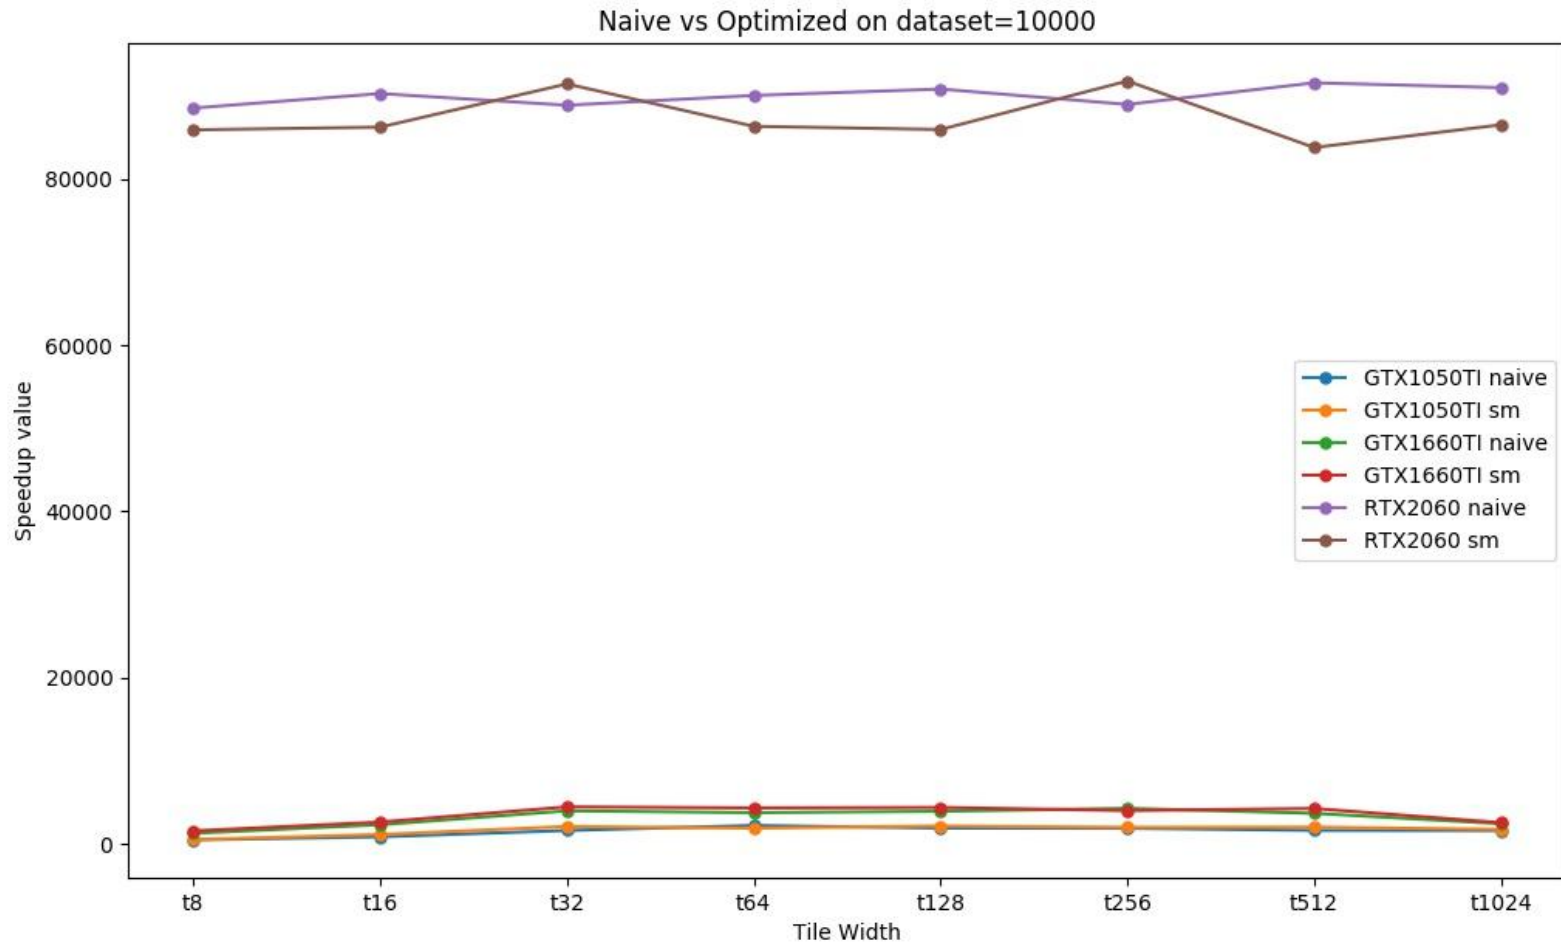
Turing

6144 MB dedicated and 1920 CUDA Cores

- **Experiments** with:

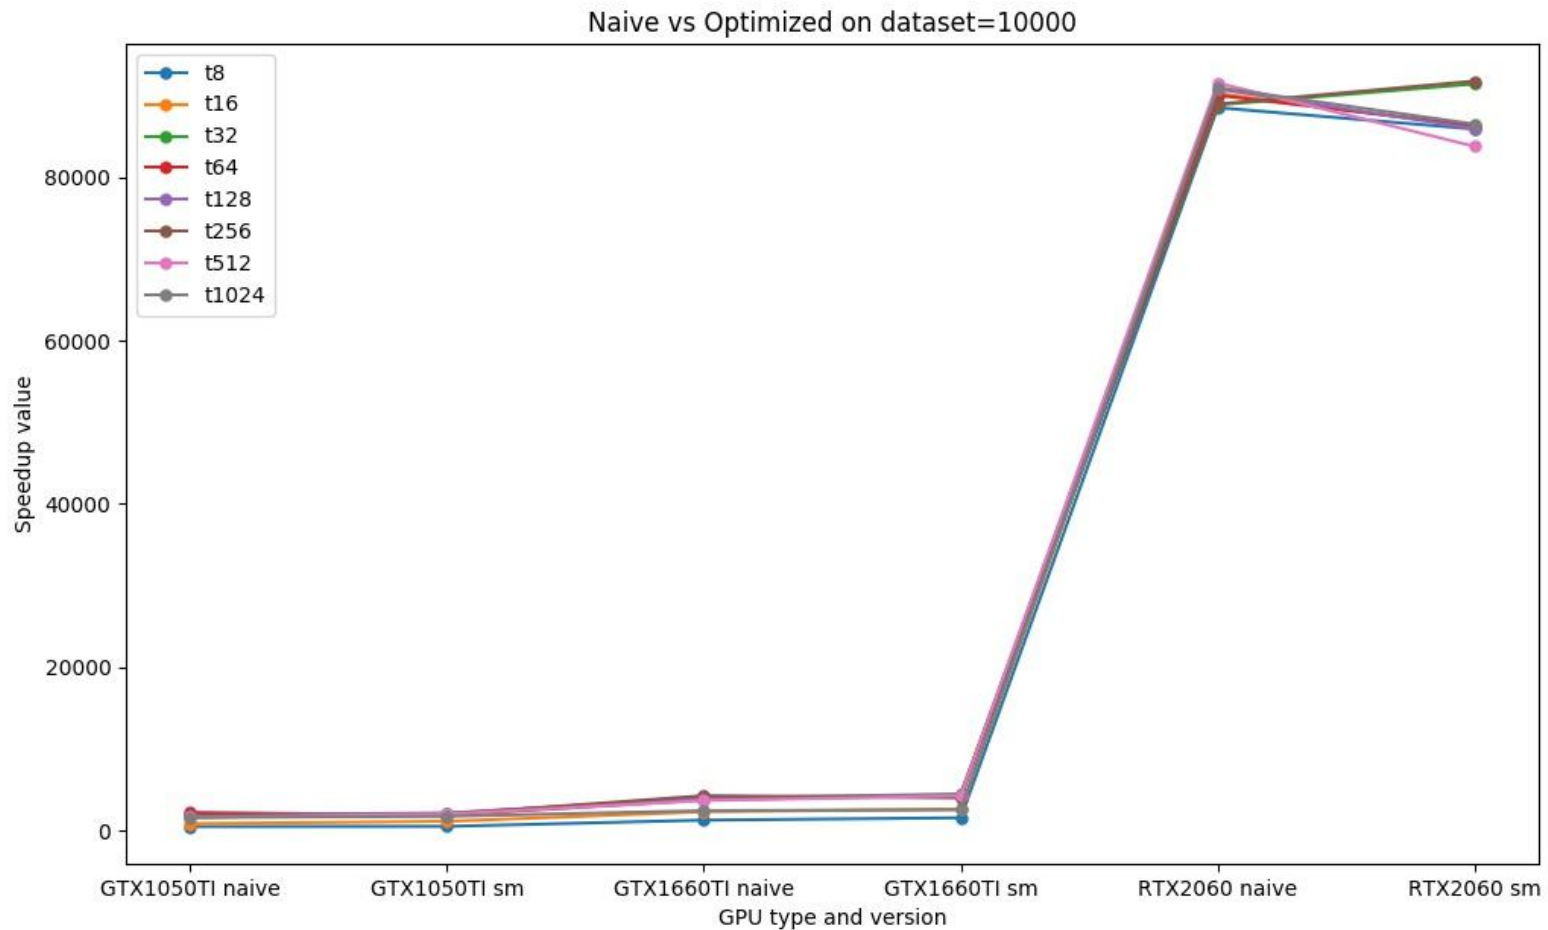
- Dataset = {10000}
- Tile Width = {8, 16, 32, 64, 128, 256, 512, 1024}

CUDA: comparison results (speedups)



RTX 2060 max speedup 91764 at TW=256

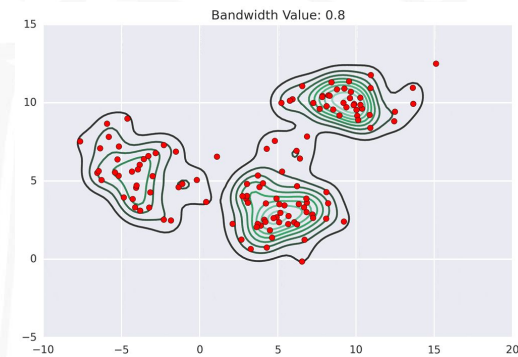
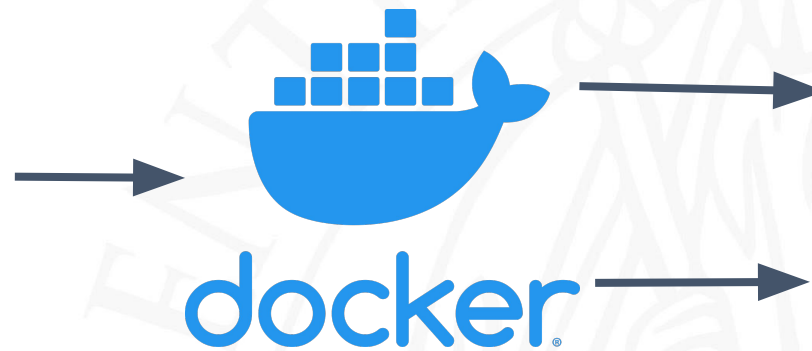
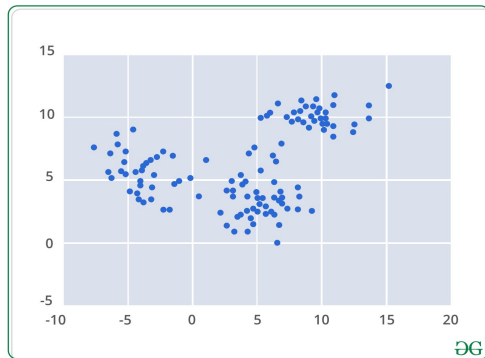
CUDA: comparison results (speedups)



RTX 2060 max speedup 91764 at TW=256

Conclusions:

Three implementations of MM: one sequential and two parallel.



<https://github.com/fabian57fabian/mean-shift-in-parallel>



UNIVERSITÀ
DEGLI STUDI
FIRENZE

Thanks for your attention

Angelo D'Amante
Fabian Greavu